

CSA0707 — COMPUTER NETWORKS

Common Assignment Report

**Analyze Protocol Behaviour and Measure Network Performance
Across Alternative Topologies Using Simulation**

Submitted By (Name)	Avinash M (192521196)
Course	CSA0707 — Computer Networks
Course Outcomes Mapped	CO3, CO4, CO5
Simulation Tool Used	Cisco Packet Tracer 8.x
Date of Submission	01 / 09 / 2026

Table of Contents

1. Problem Statement and Problem Formulation	3
2. Objective and Expected Outcomes	3
3. Requirements, Constraints and Assumptions	4
4. Application of Relevant Course Knowledge / Concepts	5
5. Design / Proposed Solution / Methodology	6
6. Algorithm / Pseudocode / Flowchart	8
7. Implementation / Source Code and Environment / Tools Used	9
8. Test Cases and Expected / Actual Results	10
9. Execution Screenshots / Outputs	10
10. Results and Validation	11
11. Analysis, Comparison, Trade-offs and Justification	12
12. Broader Considerations / SDG Relevance	13
13. Conclusion, Limitations and Possible Improvements	13
14. Individual Contribution of Group Members	14
15. References	15
16. One-Page Individual Reflection	15

(Page numbers are approximate — right-click the finished document in Word and update fields, or adjust manually, once screenshots and final content are inserted.)

1. Problem Statement and Problem Formulation

Modern campus and enterprise networks must support a wide variety of applications — file transfer, live video, web browsing, and name resolution — each of which produces traffic with different characteristics in terms of bandwidth demand, delay sensitivity, and loss tolerance. Network engineers must therefore choose a topology and a set of core protocols that keep performance acceptable even as traffic conditions change.

This assignment requires the design, simulation, and evaluation of a multi-segment campus network containing at least four interconnected segments and a minimum of forty end devices. Four alternative physical/logical topologies — star, mesh, bus, and hybrid — must be implemented for the same network requirement using an appropriate simulator (Cisco Packet Tracer, GNS3, or EVE-NG). Within this simulated environment, the behaviour of TCP, UDP, HTTP, and DNS must be analyzed under at least three controlled network conditions (for example, low traffic load, high traffic load, and increased delay / packet loss / link congestion).

The core problem to be solved is therefore two-fold: (a) determine how each protocol's performance changes as network conditions are varied, and (b) determine which topology best supports the application mix while remaining scalable, reliable, and cost-effective, using measured simulation evidence rather than assumption.

2. Objective and Expected Outcomes

Objectives:

- Design a campus network with a minimum of four interconnected segments and forty end devices supporting mixed application traffic (web, file transfer, DNS lookups, and streaming/VoIP-like UDP traffic).
- Implement the same network requirement using four alternative topologies — star, mesh, bus, and hybrid — selecting appropriate devices, transmission media, and cabling standards for each.
- Simulate the network in Cisco Packet Tracer and generate controlled traffic scenarios representing low traffic, high traffic, and increased delay/packet loss/congestion.

- Measure a minimum of five key performance metrics — latency, throughput, packet loss, jitter, and successful delivery rate/link utilization — for TCP, UDP, HTTP, and DNS traffic under each condition and each topology.
- Compare protocol and topology performance quantitatively, identify bottlenecks, and justify a final recommended topology with supporting evidence.

Expected Outcomes:

- A working, testable simulation model of all four topologies carrying the same application traffic mix.
- A quantitative performance dataset (tables/graphs) that shows how TCP, UDP, HTTP, and DNS respond differently to changing network conditions.
- A justified engineering recommendation for the topology and design choices best suited to the given application requirements.

3. Requirements, Constraints and Assumptions

3.1 Functional Requirements

- Minimum of 4 interconnected network segments (e.g., Admin, Faculty, Student Lab, Server/Data Centre segment).
- Minimum of 40 end devices distributed across the segments (PCs, laptops, servers, printers).
- Each segment must be able to reach every other segment (inter-VLAN/inter-segment routing where applicable).
- Support for HTTP (web browsing to an internal web server), DNS (name resolution for internal/external names), TCP-based file transfer (FTP/TFTP or generic TCP application), and UDP-based traffic (e.g., simple UDP client-server or VoIP-like traffic).

3.2 Non-Functional Requirements

- Scalability — the design should allow additional devices/segments to be added with minimal redesign.
- Reliability — link/device redundancy should be evaluated, especially for the mesh and hybrid designs.
- Measurability — the simulator must expose or allow calculation of latency, throughput, jitter, packet loss, and delivery rate.

3.3 Constraints

- Limited to devices and features supported by the chosen simulator (Cisco Packet Tracer).
- Simulated bandwidth/delay values are used to emulate congestion since a physical testbed is not available.
- Team must complete design, implementation, testing and documentation within the 2-day submission window.

3.4 Assumptions

- All end devices run standard TCP/IP stacks with default OS traffic-generation behaviour as modelled by the simulator.
- Background/cross-traffic used to emulate 'high traffic' and 'congestion' conditions is generated using the simulator's traffic generation or multiple concurrent PDU/ping/file-transfer requests.
- Cabling distances are within the limits specified by the relevant IEEE/TIA cabling standards used for each medium.

4. Application of Relevant Course Knowledge / Concepts

This assignment draws directly on the following course concepts:

- OSI and TCP/IP reference models — mapping HTTP and DNS to the application layer, TCP/UDP to the transport layer, and topology/media/cabling decisions to the physical and data-link layers.
- Transport layer protocol behaviour — TCP's connection-oriented, reliable, congestion-controlled delivery versus UDP's connectionless, low-overhead, loss-tolerant delivery, and how each behaves differently under delay and loss.
- Application layer protocols — HTTP's request/response behaviour over TCP and its sensitivity to round-trip time (RTT); DNS's typically UDP-based, latency-critical query/response pattern.
- Network topology design — star (centralized switching), bus (shared media, legacy), mesh (full/partial device interconnection for redundancy), and hybrid (combination of the above) topologies, along with their trade-offs in cost, fault tolerance, and scalability.
- Structured cabling and transmission media — twisted pair (Cat5e/Cat6, TIA/EIA-568), fibre optic (single-mode/multi-mode), and their bandwidth/distance characteristics.

- Network performance metrics — latency, jitter, throughput, packet loss, and utilization, and how simulation tools such as Packet Tracer can be used to approximate these under controlled conditions.

5. Design / Proposed Solution / Methodology

The same logical requirement — four segments (Admin, Faculty, Student Lab, Server/Data Centre), 40+ end devices, and support for HTTP, DNS, TCP file transfer and UDP traffic — is implemented as four alternative physical topologies. Each design uses the same IP addressing scheme and application servers so that protocol behaviour can be compared fairly across topologies.

5.1 Common Building Blocks

Segment	Approx. End Devices	Key Servers
Administration	10	—
Faculty	10	—
Student Lab	15	—
Server / Data Centre	5+	DNS Server, Web (HTTP) Server, File/TFTP Server, UDP Application Server

Total end devices: 40+ (excluding intermediary devices such as switches and routers).

5.2 Star Topology

Each segment is built around a central Layer-2/Layer-3 switch to which all end devices in that segment connect directly using Cat6 UTP (TIA/EIA-568-B, 100 m max run). The four segment switches connect to a central core switch/router, forming a 'star-of-stars'.

- Devices: 4 × access switch (24-port), 1 × core Layer-3 switch/router, 1 × router for external connectivity.
- Media/Cabling: Cat6 UTP straight-through cabling to end devices; fibre or Cat6a uplinks between access and core switches.
- Advantage: Simple to install and troubleshoot; a single device failure only affects that device.
- Limitation: The core switch is a single point of failure for inter-segment traffic.

5.3 Bus Topology

All devices within a segment are attached to a single shared backbone cable (modelled in Packet Tracer using a shared coaxial/Ethernet backbone with T-connectors and terminators, or a hub to emulate a shared-medium collision domain), and the four segment backbones are joined through a single router.

- Devices: 1 × shared backbone (hub-emulated) per segment, 1 × router joining the four segments.
- Media/Cabling: Coaxial (10BASE2-style) or hub-based UTP to emulate a shared collision domain.
- Advantage: Low cabling cost.
- Limitation: A single cable break disrupts the whole segment; performance degrades sharply under high traffic due to collisions.

5.4 Mesh Topology

Core devices (segment switches/routers) are interconnected with redundant links so that multiple paths exist between any two segments (partial mesh at the core, full mesh optionally simulated among the four core devices).

- Devices: 4 × segment switch/router interconnected with redundant fibre/Cat6a links; dynamic routing (e.g., OSPF) enabled for path redundancy.
- Media/Cabling: Fibre optic (single-mode, TIA-568 compliant) for core links; Cat6 UTP to end devices.
- Advantage: High fault tolerance — traffic can reroute around a failed link.
- Limitation: Higher device, cabling and configuration cost/complexity.

5.5 Hybrid Topology

Each segment internally uses a star topology (for ease of management), while the core interconnection between the four segments uses a partial-mesh design (for redundancy) — combining the strengths of both.

- Devices: 4 × access switch (star, per segment) + core switches interconnected in partial mesh.

- Media/Cabling: Cat6 UTP to end devices; redundant fibre/Cat6a links at the core.
- Advantage: Balances manageability, cost and fault tolerance; used as the basis for the final recommendation (see Section 11).

5.6 Common Addressing and Application Placement

A consistent private IP addressing plan (e.g., 192.168.10.0/24 – Admin, 192.168.20.0/24 – Faculty, 192.168.30.0/24 – Student Lab, 192.168.40.0/24 – Server segment) is reused across all four topology files so that only the physical/logical interconnection changes between designs, keeping the comparison fair.

6. Algorithm / Pseudocode / Flowchart

The following pseudocode describes the controlled experiment procedure applied identically to each of the four topologies.

```
BEGIN Experiment(topology T, protocol P, condition C)
1. Load simulation file for topology T (STAR | BUS | MESH | HYBRID)
2. Configure end devices, servers and IP addressing (identical across all T)
3. SET network condition C:
   IF C == LOW_TRAFFIC -> generate 1 concurrent session per protocol
   IF C == HIGH_TRAFFIC -> generate N (N>=10) concurrent sessions per protocol
   IF C == DEGRADED_LINK -> inject delay/loss on selected link(s) using
                           simulated WAN/link properties
4. FOR each protocol P in {TCP, UDP, HTTP, DNS}:
   a. Generate representative traffic for P
      (TCP: file transfer; UDP: streaming/VoIP-like traffic;
       HTTP: repeated page requests; DNS: repeated name lookups)
   b. Start packet capture / simulation logging
   c. RUN traffic for fixed duration (e.g., 5 simulated minutes)
   d. RECORD: latency, throughput, packet loss, jitter, delivery rate
   e. STOP capture; SAVE results to results_T_P_C.csv
5. REPEAT steps 3-4 for each condition C (at least 3 conditions)
6. REPEAT steps 2-5 for each topology T (STAR, BUS, MESH, HYBRID)
7. AGGREGATE all results_*.csv into master comparison table
8. COMPUTE averages / identify bottlenecks per topology-protocol-condition
9. OUTPUT comparison tables and graphs
END
```

A simplified flow of the procedure is: Select Topology → Build & Verify Connectivity → Apply Traffic Condition → Generate Protocol Traffic → Capture Metrics → Repeat for Next Protocol/Condition/Topology → Aggregate & Compare → Draw Conclusions.

7. Implementation / Source Code and Environment / Tools Used

7.1 Environment / Tools

Tool / Resource	Purpose
Cisco Packet Tracer 8.x	Primary network simulation environment for building all four topologies
Packet Tracer Simple PDU / Complex PDU tools	Generating and tracing individual test packets (ping, traceroute)
Packet Tracer Server services (HTTP, DNS, TFTP/FTP)	Hosting application services used to generate TCP/UDP/HTTP/DNS traffic
Packet Tracer real-time/simulation mode statistics	Capturing latency, delay and delivery outcomes for controlled tests
Spreadsheet software (e.g., Excel/Google Sheets)	Tabulating exported results and generating comparison graphs
GitHub	Version control and submission of source/simulation files

7.2 Implementation Steps

1. Create one .pkt project file per topology: star.pkt, bus.pkt, mesh.pkt, hybrid.pkt.
2. Place and cable devices per the design in Section 5 for each topology, using the specified media and cabling standards.
3. Configure IP addressing, VLANs (where applicable), and routing (static routes or OSPF for the mesh/hybrid designs).
4. Deploy DNS, HTTP and TFTP/FTP server services on the server segment; configure a UDP-based application (e.g., a simple UDP client/server pair) to emulate real-time traffic.
5. Verify baseline connectivity between all segments (ping/traceroute) before running controlled experiments.
6. Apply the three controlled conditions (low traffic, high traffic, degraded link) as described in Section 6, and log results for each protocol.
7. Export results and organize them into the tables in Section 9.

Note to team: Insert your actual Packet Tracer (.pkt) file names, device counts, and any configuration snippets (VLAN config, OSPF config, ACLs, etc.) your team implemented. Paste key configuration excerpts here as code blocks.

8. Test Cases and Expected / Actual Results

The following controlled conditions were applied to every topology (repeated for TCP, UDP, HTTP and DNS traffic):

Test ID	Condition	Description	Expected Behaviour
---------	-----------	-------------	--------------------

TC-1	Low Traffic	1 concurrent session per protocol, no induced delay/loss	Low latency, near-zero loss, stable throughput for all protocols
TC-2	High Traffic	10+ concurrent sessions per protocol competing for bandwidth	Increased latency/jitter; TCP throttles via congestion control, UDP throughput may drop with rising loss
TC-3	Increased Delay / Packet Loss / Congestion	Simulated WAN delay and/or induced loss on a core link	TCP retransmits and slows down; UDP/HTTP/DNS show visible loss/latency increase; DNS timeouts possible

Actual results obtained from simulation for each test case, topology and protocol are recorded in Section 9 (Results and Validation). Each test was repeated three times per topology/condition/protocol combination and the mean value reported, per the assignment requirement to repeat tests where appropriate.

9. Execution Screenshots / Outputs

Note to team: This is a placeholder section. Insert screenshots of: (1) each of the four completed topologies in Packet Tracer, (2) device/IP configuration screens, (3) Simple PDU / Complex PDU results for each protocol test, (4) simulation-mode statistics panels showing captured latency/loss, and (5) any exported CSV/graph output. Caption every figure (e.g., 'Figure 1: Star Topology — Full Network Layout').

[Insert Figure 1 — Star Topology layout]

[Insert Figure 2 — Bus Topology layout]

[Insert Figure 3 — Mesh Topology layout]

[Insert Figure 4 — Hybrid Topology layout]

[Insert Figure 5 — Example Simple PDU / Complex PDU trace for an HTTP request]

[Insert Figure 6 — Simulation-mode statistics panel showing captured metrics under TC-2 (High Traffic)]

10. Results and Validation

Note to team: The values in the tables below are illustrative placeholders showing the expected format and typical direction of results (e.g., mesh/hybrid performing better under congestion, bus degrading fastest under high traffic). Replace every figure with your team's actual measured values exported from Packet Tracer before submission — do not submit fabricated numbers as real results.

10.1 Metric Set Used

- Latency (ms) — average round-trip time per request
- Throughput (kbps) — effective data delivered per second
- Packet Loss (%) — proportion of packets not successfully delivered
- Jitter (ms) — variation in inter-packet arrival time (UDP-focused)
- Successful Delivery Rate (%) / Link Utilization (%)

10.2 Sample Results — High Traffic Condition (TC-2)

Topology	Protocol	Latency (ms)	Throughput (kbps)	Packet Loss (%)	Jitter (ms)	Delivery Rate (%)
Star	TCP	42	850	0.8	—	99.2
Star	UDP	38	780	2.1	6.4	97.9
Bus	TCP	95	420	4.5	—	95.5
Bus	UDP	88	360	7.8	14.2	92.2
Mesh	TCP	28	970	0.3	—	99.7
Mesh	UDP	25	910	1.0	3.1	99.0
Hybrid	TCP	31	930	0.5	—	99.5
Hybrid	UDP	29	870	1.3	3.8	98.7

10.3 Sample Results — HTTP and DNS (All Conditions)

Topology	Condition	HTTP Response Time (ms)	DNS Resolution Time (ms)
Star	Low Traffic	55	12
Star	High Traffic	140	35
Star	Degraded Link	310	90
Bus	Low Traffic	80	18
Bus	High Traffic	260	70
Bus	Degraded Link	520	150
Mesh	Low Traffic	40	9
Mesh	High Traffic	95	22
Mesh	Degraded Link	150	40
Hybrid	Low Traffic	45	10
Hybrid	High Traffic	105	25

Hybrid	Degraded Link	175	45
--------	---------------	-----	----

Graphs: export the above tables into bar charts (topology vs. latency/throughput per protocol) and line charts (metric vs. condition, one line per topology) using spreadsheet software, and insert them here as figures with captions (e.g., 'Figure 7: TCP Throughput by Topology under High Traffic').

11. Analysis, Comparison, Trade-offs and Justification of the Final Solution

11.1 Protocol Behaviour

- TCP consistently shows higher successful-delivery rates than UDP across all topologies because of retransmission and congestion control, at the cost of higher latency under congestion as it backs off.
- UDP shows lower latency under low traffic (no connection setup/ACK overhead) but higher packet loss and jitter under high traffic or degraded-link conditions, since it has no built-in recovery.
- HTTP response time closely tracks the underlying TCP performance of the topology, since HTTP sessions ride over TCP.
- DNS resolution time is the most sensitive to link degradation, since a single dropped UDP query typically forces a client-side timeout and retry, disproportionately increasing effective delay.

11.2 Topology Comparison

- Bus performs worst across all metrics, especially under high traffic, due to the shared collision domain — this matches its expected theoretical behaviour.
- Star performs adequately under low/medium traffic but the core switch/link becomes a bottleneck under high traffic, since all inter-segment traffic funnels through a single core path.
- Mesh delivers the best raw performance due to redundant paths and load distribution, but at higher device/cabling/configuration cost.
- Hybrid delivers performance close to mesh while keeping cost and management complexity closer to star, by combining a manageable star access layer with a redundant mesh core.

11.3 Bottlenecks Identified

- Bus: shared-medium collisions under concurrent sessions.
- Star: single core switch/uplink under high inter-segment traffic.
- Mesh/Hybrid: minor bottleneck only at the specific link where delay/loss was deliberately injected (TC-3), confirming that redundant paths successfully reroute part of the load.

11.4 Final Recommendation

Based on the measured results, the Hybrid topology is recommended as the final design: it provides near-mesh performance and fault tolerance at the network core, while retaining the simpler, more cost-effective star wiring at the access layer for each segment. This balances the assignment's performance objectives against practical cost, scalability and manageability constraints.

12. Broader Considerations / SDG Relevance

- Reliability and Resilience: The redundant links in the mesh/hybrid core reduce single points of failure, supporting dependable connectivity for critical services (e.g., DNS, web access).
- Scalability: The hybrid design allows new segments or devices to be added at the access layer without redesigning the core.
- Efficient Resource Use: Structured cabling standards (TIA/EIA-568) and appropriately sized links avoid over-provisioning while still meeting performance targets identified in testing.
- Responsible Experimentation: All traffic used for congestion/degradation testing was generated within the isolated simulation environment; no external/production network was affected.
- SDG 9 — Industry, Innovation and Infrastructure: The assignment directly practises the design of resilient, efficient digital infrastructure, aligning with SDG 9's goal of building reliable, sustainable and resilient infrastructure to support economic development and human well-being.

13. Conclusion, Limitations and Possible Improvements

13.1 Conclusion

Across all four simulated topologies, protocol behaviour matched expected theory: TCP traded latency for reliability, UDP traded reliability for speed, and both degraded under congestion — but by different amounts depending on the underlying topology. The hybrid topology offered the best balance of measured performance, cost and resilience, and is recommended as the final design for the given application requirements.

13.2 Limitations

- Packet Tracer approximates real-world link behaviour (queuing, congestion control, propagation delay) and does not perfectly reproduce a physical testbed.
- Traffic generation was limited to the concurrent-session counts and durations feasible within the simulator and the assignment time window.
- Only a fixed set of controlled conditions (three) was tested per the assignment minimum; a wider range of loss/delay values could reveal additional non-linear effects.

13.3 Possible Improvements

- Repeat experiments in GNS3/EVE-NG with real router/switch images or on a physical testbed to validate simulator results.
- Introduce QoS (traffic prioritization) for latency-sensitive UDP/DNS traffic and re-measure the same metrics.
- Automate metric collection and graphing (e.g., via scripted export) to allow a larger number of repeated trials.

14. Individual Contribution of Group Members

Name	Roll No.	Contribution
Avinash M	192521196	The most significant challenge was designing controlled experiments that isolated the effect of traffic condition (low traffic, high traffic, degraded link) from the effect of topology
—	—	—
—	—	—

Note: Update this table with the actual names, roll numbers and specific contributions of every team member — this is a mandatory component of the submission.

15. References

8. Cisco Systems, Cisco Packet Tracer — Official Documentation. Available at: <https://www.netacad.com/courses/packet-tracer>
9. Kurose, J.F. and Ross, K.W., Computer Networking: A Top-Down Approach, Pearson Education.
10. Forouzan, B.A., Data Communications and Networking, McGraw-Hill Education.
11. TIA/EIA-568 Structured Cabling Standard — Telecommunications Industry Association.
12. IETF RFC 793 — Transmission Control Protocol (TCP).
13. IETF RFC 768 — User Datagram Protocol (UDP).
14. IETF RFC 1035 — Domain Names: Implementation and Specification (DNS).
15. United Nations, Sustainable Development Goal 9 — Industry, Innovation and Infrastructure. Available at: <https://sdgs.un.org/goals/goal9>
16. [Add any additional datasets, tools or sources actually used by your team.]

16. One-Page Individual Reflection — Avinash M

Note to team: Each team member must submit their own individual reflection (300–500 words). Edit the reflection below so it genuinely reflects your own experience, decisions, challenges and learning — do not submit it unedited, and each teammate should write their own version.

Working on this assignment gave me practical experience in translating theoretical networking concepts — topology design, transport-layer behaviour, and structured cabling — into a working simulated network. One of the first design decisions I was involved in was choosing appropriate devices and media for each of the four topologies; this required revisiting cabling standards (TIA/EIA-568) and thinking carefully about trade-offs between cost, redundancy and manageability rather than simply following a textbook diagram.

The most significant challenge was designing controlled experiments that isolated the effect of traffic condition (low traffic, high traffic, degraded link) from the effect of topology, so that the two variables could be compared fairly. This meant keeping the IP addressing scheme, server placement and application mix identical across all four topology files, and only changing the interconnection and cabling. Generating realistic 'high traffic' conditions inside a simulator, and interpreting the resulting latency, jitter and loss figures correctly, deepened my understanding how TCP's congestion control and UDP's lack of retransmission produce very different behaviour under the same stress.

Comparing the measured results across topologies confirmed, with evidence rather than assumption, why redundant (mesh/hybrid) designs are preferred in real enterprise networks despite their higher cost — a conclusion that is easy to state from theory but became much clearer after seeing bus and star topologies visibly degrade under the same load that mesh and hybrid handled with far less impact.

This assignment directly supports CO3 (comparing TCP/UDP/HTTP/DNS performance under varying conditions), CO4 (designing star, mesh, bus and hybrid topologies with appropriate devices, media and cabling), and CO5 (measuring performance using key metrics through simulation). It also connects to SDG 9 by practising the design of resilient, efficient network infrastructure. Overall, the exercise strengthened both my technical skills in network simulation and my ability to justify engineering decisions using measured evidence rather than intuition alone.

